

Homework 2^{†‡}

Problem 1: Linear Regression

- (a) I loaded the data from `data/curve80.data` into the `curve` variable. Then the variable `X` stores the data and `Y` stores the regressed values. I then split the data into a 75-25 partition and store it in the individual variables. Shown in Listing 1.

Listing 1: Fetching the `data/curve80.data` and splitting it with a 75-25 partition.

```
1 % Fetching the dataset and separating it into X and Y.
2 curve=load('data/curve80.txt'); % load the text file
3 y = curve(:,end); % target value is last column
4 X = curve(:,1); % features are other columns
5 whos
6
7 % Part (a)
8 % Splitting the data into 75-25 split.
9 [Xtr, Xte, Ytr, Yte] = splitData(X,y, .75);
10
11 >> whos
12 Name Size Bytes Class Attributes
13
14 X 80x1 640 double
15 Xte 20x1 160 double
16 Xtr 60x1 480 double
17 Yte 20x1 160 double
18 Ytr 60x1 480 double
19 curve 80x2 1280 double
20 y 80x1 640 double
```

- (b) After loading the data, we use the `linearRegress` class's `predict()` to get the prediction of the model. The code is given in Listing 2. The plot can be seen at Figure 1.

Listing 2: Getting a prediction for the model built on the Training data.

```
1 %% Part (b)
2 % Training the linear regression with Xtr as an input
3 lr = linearRegress( Xtr, Ytr );
4 xs = [0:.05:10]';
5 ys = predict(lr, xs);
6 ys1 = predict(lr, Xte);
7
8 % Plotting the training data and the predicted function in the same plot.
9 f=figure;
10 scatter(Xtr, Ytr, 'filled');
11 hold on;
12 plot(xs, ys);
13 hold off;
14 saveas(f, 'lr.jpg', 'jpg');
15
16 % MSE for both Training and Test data
17 mse(lr, Xtr, Ytr)
```

*Website: <http://sli.ics.uci.edu/Classes/2015W-273a>

[†]Questions available at <http://sli.ics.uci.edu/Classes/2015W-273a?action=download&upname=HW2.pdf>

[‡]All the figures and listing numbers are auto-referred.




Figure 1: Training data and the learned function in Listing 2

```

18 % ans = 1.1277
19 mse(lr, Xte, Yte)
20 % ans = 2.2423

```

- (c) The higher degree polynomials are created as per the Listing 3. The higher-degree polynomials are created by the `fpoly()` function provided. Using the new dataset, we rescale it and train the linear regression model using `linearRegress` class.

Listing 3: Getting a prediction for the model built on the Training data.

```

1 f = figure;
2 i = 1;
3 errorsTr = zeros(1, 6);
4 errorsTe = zeros(1, 6);
5 d=[1,3,5,7,10,18];
6 for degree=[1,3,5,7,10,18];
7
8     XtrP = fpoly(Xtr, degree, false); % create poly features up to given ...
          degree; no "1" feature.
9     [XtrP, M,S] = rescale(XtrP); % it's often a good idea to scale the ...
          features
10    lr = linearRegress( XtrP, Ytr ); % create and train model
11
12    % XS vs YS plot.
13    xs = [0:.05:10]';
14    xsP = rescale( fpoly(xs,degree,false), M,S);
15    ysP = predict( lr, xsP );
16
17    % Scatter plot and the learned function on the data.
18    subplot(2,3,i);
19    title(strcat({'degree= '},num2str(degree)))
20    hold on;
21    scatter(Xtr, Ytr, 'filled');
22    ax = axis;
23    hold on;
24    plot(xs, ysP);
25    axis(ax);
26

```

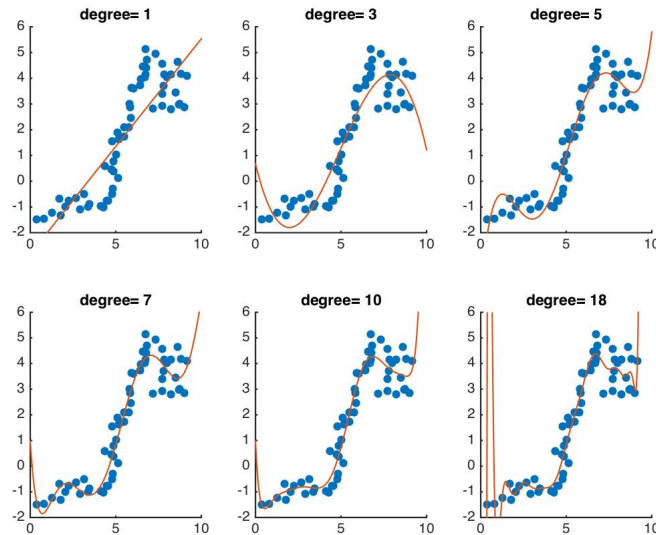


Figure 2: Plots for different degree polynomials and the learned functions

```

27 % For getting the right degree to get the best learning function.
28 XteP = rescale(fpoly(Xte,degree,false), M,S);
29 YteP = predict(lr,XteP);
30
31 % MSE for both Training and Test data
32 errorsTr(1, i) = mse(lr, XtrP, Ytr)
33 errorsTe(1, i) = mse(lr, XteP, Yte)
34 i= i+1;
35 % now, apply the same polynomial expansion & scaling transformation ...
    to Xtest:
36 %
37 end;
38 hold off;

```

The trained functions and the scatter for various degree values can be seen in the visual representation seen in Figure 2. To get to a better approximation of the right degree, we learned the performance of the learner with the test data and plotted it, as can be seen in Figure 3. From this observation, I can see that for **degree = 10**, the test error was the least.

Problem 2: Cross Validation

For each degree, I use cross validation within the loop iterating on all the degrees. Before I apply the cross-validation on this data and change its degree as well as rescale it, as seen in Listing 4. After that, I save the errors into an array and use that to plot, which can be seen in 4(a). From this plot, it can be seen that for **degree = 5**, the error is minimum. When we compare it with the error on the actual test data, the difference shows up. The actual test data fits better for degree = 10, whereas cross validated data fits better with degree = 5. This plot shows that error for cross-validated model follows similar pattern to a test data, thus we can build models on data which would generalize better, without the knowledge of the actual test data.

Listing 4: Training linear regression with cross-validation.

```

1 i = 1;
2 nFolds = 5;

```





Figure 3: Train and Test Errors.


```

3 d=[1,3,5,7,10,18];
4 for degree=[1,3,5,7,10,18];
5     % Degrees and scaling of the data
6     XtrP = fpoly(Xtr, degree, false); % create poly features up to given ...
        degree; no "1" feature.
7     [XtrP, M,S] = rescale(XtrP); % it's often a good idea to scale the features
8     for iFold = 1:nFolds;
9         [Xti,Xvi,Yti,Yvi] = crossValidate(XtrP,Ytr,nFolds,iFold);
10        % take ith data block as v learner = linearRegress(...)
11        lr = linearRegress( Xti, Yti ); % create and train model
12        % TODO: train on Xti, Yti , the data for this fold J(iFold) = ...
13
14        % TODO: now compute the MSE on Xvi, Yvi and save it
15        J(iFold) = mse(lr, Xvi, Yvi);
16    end;
17    % the overall estimated validation performance is the average of the ...
        performance on ea
18    errors(i) = mean(J);
19    i = i + 1;
20 end;

```



(a) Training Errors for various degrees when the data is cross-validated



(b) Train Errors on Cross-validated error as well as the Error on the Actual Test data

Figure 4: Running Various Experiments. The three bars are for Brute Force, I-Consistency and I-Consistency with MRV search. All the times given are in Seconds