

CS273a Project Description
Intro to Machine Learning: Fall 2012
Due: Friday December 14th, 2012

For your course project, you will explore data mining and prediction “in the wild”, in a real life data set and compared against the performance of teams from around the world.

We will use a data set from a past Knowledge Discovery in Data (KDD) Cup, a yearly competition in machine learning and data mining associated with the KDD conference. In particular, we will use the 2004 Competition’s Particle Physics data set.

The challenge is described in full on the webpage:

<http://osmot.cs.cornell.edu/kddcup/>

From the webpage:

The goal in this task is to learn a classification rule that differentiates between two types of particles generated in high energy collider experiments. It is a binary classification problem with 78 attributes. The training set has 50,000 examples, and the test set is of size 100,000.

Although the original challenge was to compete with respect to four metrics (accuracy, ROC, cross-entropy, and “q-score”), you may pick any one (or more than one) to compete with.

Step 1: Form teams.

You may work in teams of 2-3 students. Choose a name for your team; to make it recognizable to me on the leaderboard, please make it "cs273-[your-team-name]".

There is no barrier to collaboration on this project, so feel free to talk to other teams about what they are doing, how they are doing it, and how well it works (and you can see this last point on the leaderboard as well). However, please make sure that your own entries are developed by your team members,

Step 2: Download the data.

To download the data, you must be registered on the site; see:

<http://osmot.cs.cornell.edu/kddcup/datasets.html>

To load the data directly into Matlab, you will have to replace the ‘?’ in the test set files with some numeric value.

Note that some data have missing features, denoted with a 999 or 9999 value (see webpage and below). You may remove these features, ignore them, set them to zero, set them to the mean or median of the non-missing values, or impute them in any other way that you choose.

Step 3: Build your learners

Use the techniques we have developed so far in class to construct predictive models for your target(s). You may use the Matlab code provided from class; online code packages such as Weka, PMTK, or others; or write your own.

I suggest that you try several different models, such as nearest neighbor methods, decision trees, linear classifiers (logistic regression, support vector machines, etc.), naive Bayes classifiers, and/or boosted classifiers (decision stumps, etc.). Each member of your team may try one or two models, and can explore fitting them to the data and assessing their performance using validation or cross-validation.

After learning several models, you may also want to combine them using bagging or by learning a predictor from the classifier outputs.

Some areas for possible exploration:

- (a) ROC scoring: One official scoring works by computing an ROC curve. While any binary predictor can be converted into a trivial ROC curve, this may not be the most advantageous approach. “Soft” scores (for example, unthresholded predictor values) will be converted into an ROC in a more fine-grain and possibly better way.
- (b) Q-scoring: Q-scoring is some non-standard score that the physics domain experts care about. You could redesign an algorithm to explicitly attempt to optimize Q-scoring.
- (c) Missing data: some features have missing values. These can be dealt with by simply ignoring those features; setting their values to zero; setting their values to the median or mean of the data with non-missing values; building an unsupervised model to impute their values; or many other approaches.
- (d) Feature selection or creation: the data have about 78 features. You may want to use feature selection to determine which features are useful, and reduce the number of features to avoid overfitting; on the other hand, you may want to generate new features if you feel that a model is actually under-fitting. See my functions like `fProject` and `fKitchenSink` for ideas.

Step 4: Evaluate; go to Step 3.

Output your predictions to a set of files as described in the instructions page:

<http://osmot.cs.cornell.edu/kddcup/submission.html>

You can check the leaderboard to see your “test” performance:

<http://osmot.cs.cornell.edu/cgi-bin/newtable.pl?prob=phy>

Please limit yourself to at most a few (less than 5) submissions per day to avoid over-loading their servers.

Step 5: Write it up

Your team will produce a single write-up document, approximately **6 pages** long, describing the problem you chose to tackle (which loss, etc.) and the methods you used to address it, including which model(s) you tried, how you trained them, how you selected any parameters they might require, and how they performed in on the test data. Consider including tables of performance of different approaches, or plots of performance used to perform model selection (i.e., parameters that control complexity).

Within your document, please try to describe to the best of your ability who was responsible for which aspects (which learners, etc.), and how the team as a whole put the ideas together.

You are free to collaborate with other teams, *including* sharing ideas and even code, but please document where your predictions came from. For example, for any code you use, please say who wrote the code and how it was applied (who determined the parameter settings and how, etc.) Collaboration is particularly true for learning ensembles of predictors: your teams may each supply a set of predictors, and then collaborate to learn an ensemble from the set.