

+

Machine Learning and Data Mining

Linear classification


Prof. Alexander Ihler




Supervised learning

- Notation

- Features x
- Targets y
- Predictions \hat{y}
- Parameters θ



Linear regression



“Predictor”:

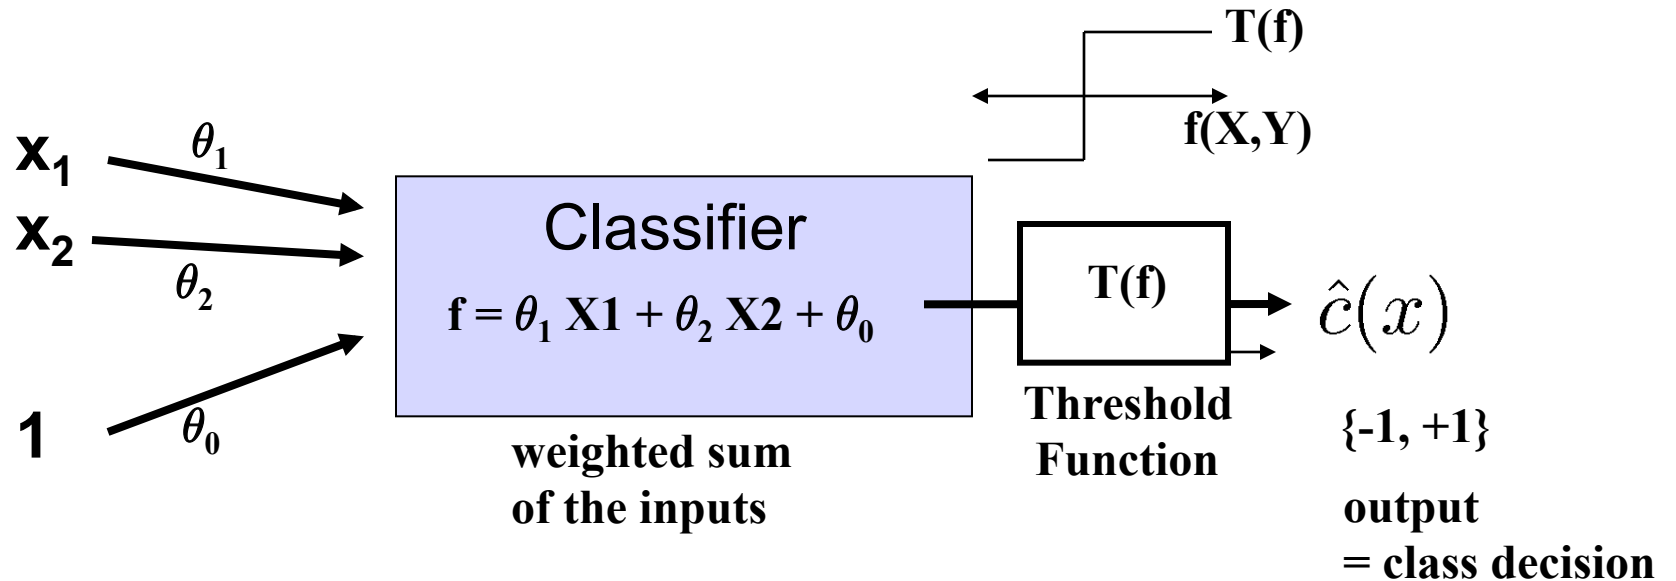
Evaluate line:

$$r = \theta_0 + \theta_1 x_1$$



return r

- Contrast with classification
 - Classify: predict discrete-valued target y

Perceptron Classifier (2 features)



Visualizing for one feature “x”:



Perceptrons

- Perceptron = a linear classifier
 - The parameters θ are sometimes called weights (“w”)
 - real-valued constants (can be positive or negative)
 - Define an additional constant input “1”
- A perceptron calculates 2 quantities:
 - 1. A weighted sum of the input features
 - 2. This sum is then thresholded by the $T(\cdot)$ function
- Perceptron: a simple artificial model of human neurons
 - weights = “synapses”
 - threshold = “neuron firing”

Notation

- Inputs:
 - $x_0, x_1, x_2, \dots, x_d$,
 - $x_1, x_2, \dots, x_{d-1}, x_d$ are the values of the d features
 - $x_0 = 1$ (a constant input)
 - $\underline{x} = (x_0, x_1, x_2, \dots, x_d)$: feature vector (row vector)
- Weights (parameters):
 - $\theta_0, \theta_1, \theta_2, \dots, \theta_d$,
 - we have $d+1$ weights
 - one for each feature + one for the constant
 - $\underline{\theta} = (\theta_0, \theta_1, \theta_2, \dots, \theta_d)$: parameter vector (row vector)
- Linear response
 - $\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_d x_d = \underline{\theta} \cdot \underline{x}'$ then threshold

(Matlab) `>> f = th*x'; f = sum(th.*x); yhat = sign(f);`

Perceptron Decision Boundary

- The perceptron is defined by the decision algorithm:

$$o(x_1, x_2, \dots, x_d, x_{d+1}) \begin{cases} = 1 & (\text{if } \underline{\theta} \cdot \underline{x}' > 0) \\ = -1 & (\text{otherwise}) \end{cases}$$


- The perceptron represents a hyperplane decision surface in d-dimensional space
 - A line in 2D, a plane in 3D, etc.
- The equation of the hyperplane is given by

$$\underline{\theta} \cdot \underline{x}' = 0$$

This defines the set of points that are on the boundary.


Example, Linear Decision Boundary

$$\begin{aligned}\underline{\theta} &= (\theta_1, \theta_2, \theta_0) \\ &= (.5, -.5, 0)\end{aligned}$$



Example, Linear Decision Boundary

$$\begin{aligned}\underline{\theta} &= (\theta_1, \theta_2, \theta_0) \\ &= (.5, -.5, 0)\end{aligned}$$




Example, Linear Decision Boundary

$$\begin{aligned}\underline{\theta} &= (\theta_1, \theta_2, \theta_0) \\ &= (.5, -.5, 0)\end{aligned}$$



$$\underline{\theta} \cdot \underline{x}' < 0$$

$\Rightarrow x_1 < x_2$
(this is the
equation for
decision
region -1)




Separability

- A data set is separable by a learner if
 - There is some instance of that learner that correctly predicts all the data points
- Linearly separable data
 - Can separate the two classes using a straight line in feature space
 - in 2 dimensions the decision boundary is a straight line




Class overlap


- Classes may not be well-separated
- Same observation values possible under both classes
 - High vs low risk; features {age, income}
 - Benign/malignant cells look similar
 - ...
- Common in practice
- May not be able to perfectly distinguish between classes
 - Maybe with more features?
 - Maybe with more complex classifier?
- Otherwise, may have to accept some errors



Another example

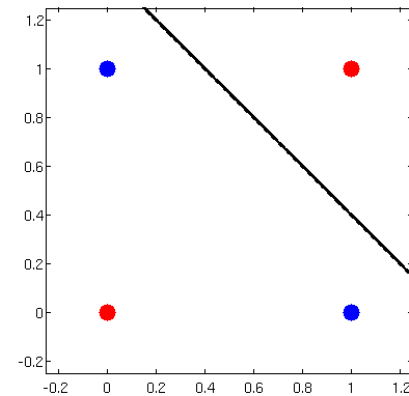



Non-linear decision boundary



Representational Power of Perceptrons

- What mappings can a perceptron represent perfectly?
 - A perceptron is a linear classifier
 - thus it can represent any mapping that is linearly separable
 - some Boolean functions like AND (on left)
 - but not Boolean functions like XOR (on right)



Adding features

- Linear classifier can't learn some functions

1D example:


$$y = T(b x + c)$$



Not linearly separable

Add quadratic features

$$y = T(a x^2 + b x + c)$$



Linearly separable in new features...

Adding features

- Linear classifier can't learn some functions

1D example:


$$y = T(bx + c)$$



Not linearly separable

Quadratic features, visualized in original feature space:

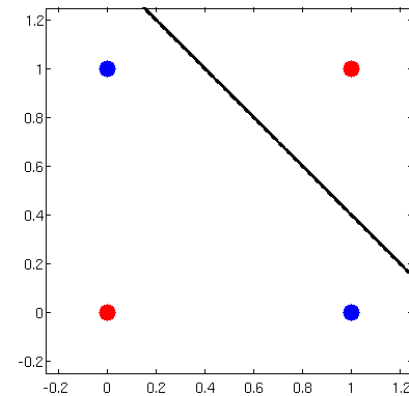

$$y = T(ax^2 + bx + c)$$



More complex decision boundary: $ax^2+bx+c = 0$

Representational Power of Perceptrons



- What mappings can a perceptron represent perfectly?
 - A perceptron is a linear classifier
 - thus it can represent any mapping that is linearly separable
 - some Boolean functions like AND (on left)
 - but not Boolean functions like XOR (on right)



What kinds of functions would we need to learn the data on the right?

Representational Power of Perceptrons

- What mappings can a perceptron represent perfectly?
 - A perceptron is a linear classifier
 - thus it can represent any mapping that is linearly separable
 - some Boolean functions like AND (on left)
 - but not Boolean functions like XOR (on right)




What kinds of functions would we need to learn the data on the right?

Ellipsoidal decision boundary: $a(x_1 - b)^2 + c(x_2 - d)^2 = 0$

Effect of dimensionality

- Data are increasingly separable in high dimension – is this a good thing?
- “Good”
 - Separation is easier in higher dimensions (for fixed N)
 - Increase the number of features, and even a linear classifier will eventually be able to separate all the training examples!
- “Bad”
 - Remember training vs. test error? Remember overfitting?
 - Increasingly complex decision boundaries can eventually get all the training data right, but it doesn't necessarily bode well for test data...



+

Machine Learning and Data Mining

Linear classification: Learning

Prof. Alexander Ihler



Learning the Classifier Parameters

- Learning from Training Data:
 - training data = labeled feature vectors
 - Find parameter values that predict well (low error)
 - error is estimated on the training data
 - “true” error will be on future test data
- Define an objective function $J(\underline{\theta})$:
 - Classifier accuracy (for a given set of weights $\underline{\theta}$ and labeled data)
- Maximize this objective function (or, minimize error)
 - An optimization or search problem over the vector $(\theta_1, \theta_2, \theta_0)$

Learning the Weights from Data

An Example of a Training Data Set

Example	x_1	x_2	x_d	true class label, y
$\underline{x}(1)$	3.4	-1.2	7.1	1
$\underline{x}(2)$	4.1	-3.1	4.6	-1
$\underline{x}(3)$	5.7	-1.0	6.2	-1
$\underline{x}(4)$	2.2	4.1	5.0	1
$\underline{x}(n)$	1.2	4.3	6.1	1


Training a linear classifier

- How should we measure error?
 - Natural measure = “fraction we get wrong” (error rate)
$$\text{err}(\underline{\theta}) = 1/N \sum \delta(\hat{y}(i) \neq y(i))$$
where $\delta(\hat{y}(i) \neq y(i)) = 0$ if $\hat{y}(i) = y(i)$, and 1 otherwise

(Matlab) `>> yh = sign(th*X'); err = mean(y ~= yh);`

- But, hard to train via gradient descent
 - Not continuous
 - As decision boundary moves, errors change abruptly


1D example:




$$\begin{aligned} T(f) &= -1 \quad \text{if } f < 0 \\ T(f) &= +1 \quad \text{if } f > 0 \end{aligned}$$

Linear regression?

- Simple option: set θ using linear regression



- In practice, this often doesn't work so well...
 - Consider adding a distant but “easy” point
 - MSE distorts the solution



Perceptron algorithm

- Perceptron algorithm: an SGD-like algorithm

While (~done)

For each data point j :

$\hat{y}(j) = T(\underline{\theta} * \underline{x}(j))$: predict output for data point j

$\underline{\theta} \leftarrow \underline{\theta} + \alpha (y(j) - \hat{y}(j)) \underline{x}(j)$: “gradient-like” step

- Compare to linear regression + MSE cost

- Identical update to SGD for MSE except error uses thresholded $\hat{y}(j)$ instead of linear response $\underline{\theta} \cdot \underline{x}'$ so:

- (1) For correct predictions, $y(j) - \hat{y}(j) = 0$

- (2) For incorrect predictions, $y(j) - \hat{y}(j) = \pm 1$

“adaptive” linear regression: correct predictions stop contributing

Perceptron algorithm


- Perceptron algorithm: an SGD-like algorithm

While (~done)



For each data point j :

$\hat{y}(j) = T(\underline{\theta} * \underline{x}(j))$: predict output for data point j

$\underline{\theta} \leftarrow \underline{\theta} + \alpha (y(j) - \hat{y}(j)) \underline{x}(j)$: “gradient-like” step



$y(j)$
predicted
incorrectly:
update
weights



Perceptron algorithm


- Perceptron algorithm: an SGD-like algorithm

While (~done)


For each data point j :

$\hat{y}(j) = T(\underline{\theta} * \underline{x}(j))$: predict output for data point j

$\underline{\theta} \leftarrow \underline{\theta} + \alpha (y(j) - \hat{y}(j)) \underline{x}(j)$: “gradient-like” step



$y(j)$
predicted
correctly:
no update



Perceptron algorithm

- Perceptron algorithm: an SGD-like algorithm


While (~done)

For each data point j :


$\hat{y}(j) = T(\underline{\theta} * \underline{x}(j))$: predict output for data point j

$\underline{\theta} \leftarrow \underline{\theta} + \alpha (y(j) - \hat{y}(j)) \underline{x}(j)$: “gradient-like” step

(Converges if data are linearly separable)




$y(j)$
predicted
correctly:
no update




Surrogate loss functions

- Another solution: use a “smooth” loss
 - e.g., approximate the threshold function



- Usually some smooth function of distance
 - Example: “sigmoid”, looks like an “S”




- Now, measure e.g. MSE

$$J(\underline{\theta}) = \frac{1}{m} \sum_j \left(\sigma(f(x^{(i)})) - y^{(i)} \right)^2$$


Class $y = \{0, 1\} \dots$

- Far from the decision boundary: $|f(\cdot)|$ large, small error
- Nearby the boundary: $|f(\cdot)|$ near $1/2$, larger error

1D example:





Classification error = MSE = $2/9$



MSE = $(0^2 + 1^2 + .2^2 + .25^2 + .05^2 + \dots)/9$

Beyond misclassification rate



- Which decision boundary is “better”?
 - Both have zero training error (perfect training accuracy)
 - But, one of them seems intuitively better...



- Side benefit of “smoothed” error function
 - Encourages data to be far from the decision boundary
 - See more examples of this principle later...



Training the Classifier

- Once we have a smooth measure of quality, we can find the “best” settings for the parameters of $f(X1,X2) = a*X1 + b*X2 + c$
- Example: 2D feature space \Leftrightarrow parameter space




Training the Classifier

- Once we have a smooth measure of quality, we can find the “best” settings for the parameters of $f(X1, X2) = a * X1 + b * X2 + c$
- Example: 2D feature space \Leftrightarrow parameter space




Training the Classifier

- Once we have a smooth measure of quality, we can find the “best” settings for the parameters of $f(X1, X2) = a \cdot X1 + b \cdot X2 + c$
- Finding the minimum loss $J(.)$ in parameter space...



- $[a \ b \ c] = ?$




$$[\arctan(a/b), c] = [-\pi/4, 1]$$

Finding the Best MSE

- As in linear regression, this is now just optimization
- Methods:
 - Gradient descent
 - Improve loss by small changes in parameters (“small” = learning rate)
 - Or, substitute your favorite optimization algorithm...
 - Coordinate descent
 - Stochastic search
 - Genetic algorithms

Gradient Descent



Gradient Equations

- MSE (note, depends on function $\sigma(\cdot)$)

$$J(\underline{\theta} = [a, b, c]) = \frac{1}{N} \sum_i (\sigma(ax_1^{(i)} + bx_2^{(i)} + c) - y^{(i)})^2$$

- What's the derivative with respect to one of the parameters?

$$\frac{\partial J}{\partial a} = \frac{1}{N} \sum_i 2(\sigma(\theta \cdot x^{(i)}) - y^{(i)}) \frac{\partial \sigma(\theta \cdot x^{(i)})}{\partial a} x_1^{(i)}$$

Error between class
and prediction

Sensitivity of prediction to
changes in parameter "a"

- Similar for parameters b, c [replace x_1 with x_2 or 1 (constant)]

Saturating Functions

- Many possible “saturating” functions
- “Logistic” sigmoid (scaled for range [0,1]) is

$$\sigma(z) = 1 / (1 + \exp(-z))$$

- Derivative is

$$\partial\sigma(z) = \sigma(z) (1-\sigma(z))$$

(to predict: threshold at $\frac{1}{2}$)

- Matlab Implementation:

```
function s = sig(z)
% value of [0,1] sigmoid
s = 1 ./ (1+exp(-z));
```

```
function ds = dsig(x)
% derivative of (scaled) sigmoid
ds = sig(z) .* (1-sig(z));
```

For range [-1 , +1]:

$$\rho(z) = 2 \sigma(z) - 1$$

$$\partial\rho(z) = 2 \sigma(z) (1-\sigma(z))$$

To predict: threshold at zero

Logistic regression

- Interpret $\sigma(\underline{\theta} \cdot \mathbf{x}')$ as a probability that $y = 1$
- Use a negative log-likelihood loss function
 - If $y = 1$, cost is $-\log \Pr[y=1] = -\log \sigma(\underline{\theta} \cdot \mathbf{x}')$
 - If $y = 0$, cost is $-\log \Pr[y=0] = -\log (1 - \sigma(\underline{\theta} \cdot \mathbf{x}'))$

- Can write this succinctly:

$$J(\underline{\theta}) = -\frac{1}{m} \sum_i \underbrace{y^{(i)} \log \sigma(\underline{\theta} \cdot \mathbf{x}^{(i)})}_{\text{Nonzero only if } y=1} + \underbrace{(1-y^{(i)}) \log (1-\sigma(\underline{\theta} \cdot \mathbf{x}^{(i)}))}_{\text{Nonzero only if } y=0}$$

Logistic regression


- Interpret $\sigma(\underline{\theta} \cdot \mathbf{x}')$ as a probability that $y = 1$
- Use a negative log-likelihood loss function
 - If $y = 1$, cost is $-\log \Pr[y=1] = -\log \sigma(\underline{\theta} \cdot \mathbf{x}')$
 - If $y = 0$, cost is $-\log \Pr[y=0] = -\log (1 - \sigma(\underline{\theta} \cdot \mathbf{x}'))$

- Can write this succinctly:


$$J(\underline{\theta}) = -\frac{1}{m} \sum_i y^{(i)} \log \sigma(\underline{\theta} \cdot \mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - \sigma(\underline{\theta} \cdot \mathbf{x}^{(i)}))$$

- Convex! Otherwise similar: optimize $J(\theta)$ via ...

1D example:



Classification error = MSE = 2/9



NLL = $-(\log(.99) + \log(.97) + \dots)/9$

Gradient Equations

- Logistic neg-log likelihood loss:


$$J(\underline{\theta}) = -\frac{1}{m} \sum_i y^{(i)} \log \sigma(\theta \cdot x^{(i)}) + (1 - y^{(i)}) \log(1 - \sigma(\theta \cdot x^{(i)}))$$

- What's the derivative with respect to one of the parameters?

$$\begin{aligned} \frac{\partial J}{\partial a} &= -\frac{1}{m} \sum_i y^{(i)} \frac{1}{\sigma(\theta \cdot x^{(i)})} \partial \sigma(\theta \cdot x^{(i)}) x_1^{(i)} + (1 - y^{(i)}) \dots \\ &= -\frac{1}{m} \sum_i y^{(i)} (1 - \sigma(\theta \cdot x^{(i)})) x_1^{(i)} - (1 - y^{(i)}) \dots \end{aligned}$$

Surrogate loss functions


- Replace 0/1 loss $\Delta_i(\theta) = \delta(T(\theta x^{(i)}) \neq y^{(i)})$ with something easier:



0/1 Loss


- Logistic MSE

$$J_i(\theta) = 4(\sigma(\theta x^{(i)}) - y^{(i)})^2$$



- Logistic Neg Log Likelihood

$$J_i(\theta) = -y^{(i)} \log \sigma(\theta \cdot x^{(i)}) + \dots$$



Summary

- Linear classifier \Leftrightarrow perceptron
- Visualizing the decision boundary
- Measuring quality of a decision boundary
 - Logistic sigmoid + MSE criterion
 - Logistic Regression
- Learning the weights of a linear classifier from data
 - Reduces to an optimization problem
 - Perceptron algorithm
 - For MSE or Logistic NLL, we can do gradient descent
 - Gradient equations & update rules
- Extending features and separability