

+

Machine Learning and Data Mining


Clustering

Prof. Alexander Ihler
Fall 2012




Unsupervised learning

- Supervised learning
 - Predict target value (“y”) given features (“x”)
- Unsupervised learning
 - Understand patterns of data (just “x”)
 - Useful for many reasons
 - Data mining (“explain”)
 - Missing data values (“impute”)
 - Representation (feature generation or selection)




Unsupervised learning

- Supervised learning
 - Predict target value (“y”) given features (“x”)
- Unsupervised learning
 - Understand patterns of data (just “x”)
 - Useful for many reasons
 - Data mining (“explain”)
 - Missing data values (“impute”)
 - Representation (feature generation or selection)
- One example: *clustering*




Clustering and Data Compression

- Clustering is related to vector quantization
 - Dictionary of vectors (the cluster centers)
 - Each original value represented using a dictionary index
 - Each center “claims” a nearby region (Voronoi region)




Hierarchical Agglomerative Clustering

Initially, every datum is a cluster





- Another simple clustering alg
- Define a distance between clusters (return to this)
- Initialize: every example is a cluster
- Iterate:
 - Compute distances between all clusters (store for efficiency)
 - Merge two closest clusters
- Save both clustering and *sequence* of cluster ops
- “Dendrogram”

Iteration 1





Iteration 2



Iteration 3



- Builds up a sequence of clusters (“hierarchical”)



- Algorithm complexity $O(N^2)$
(Why?)


In matlab: “linkage” function (stats toolbox)

Dendrogram




Cluster Distances

$$D_{\min}(C_i, C_j) = \min_{x \in C_i, y \in C_j} \|x - y\|^2$$



produces minimal spanning tree.


$$D_{\max}(C_i, C_j) = \max_{x \in C_i, y \in C_j} \|x - y\|^2$$



avoids elongated clusters.

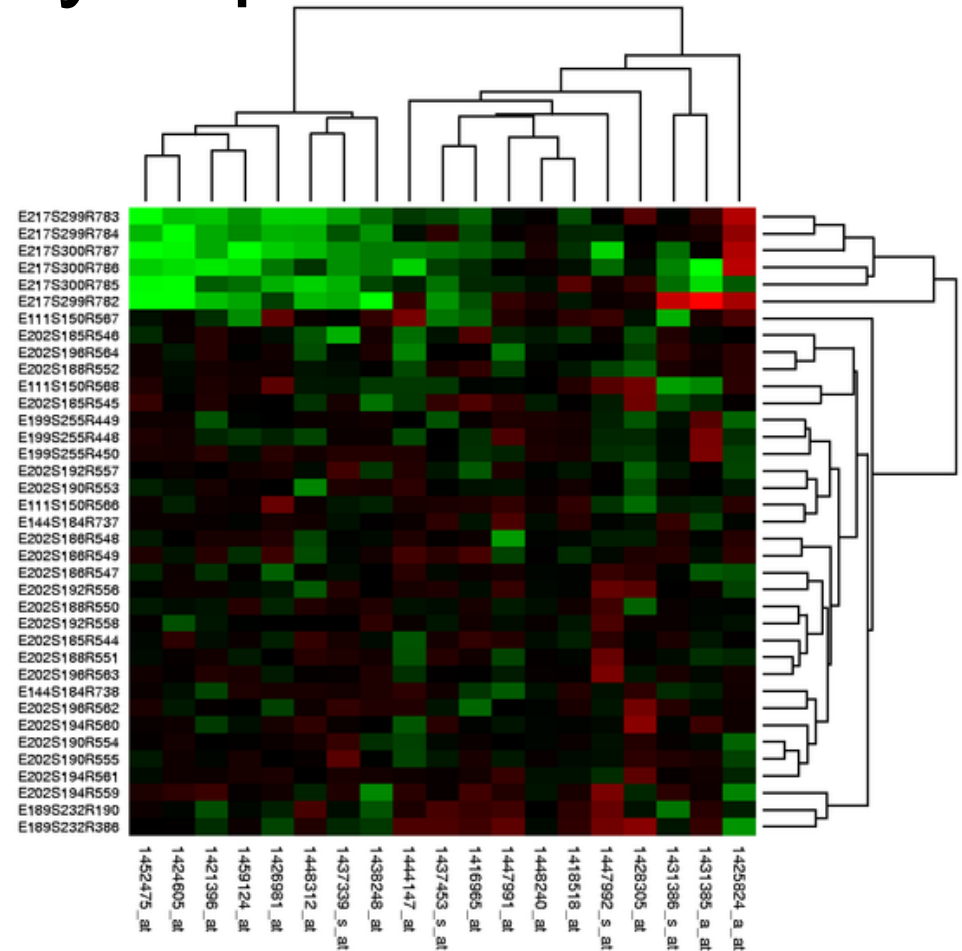
$$D_{\text{avg}}(C_i, C_j) = \frac{1}{|C_i||C_j|} \sum_{x \in C_i, y \in C_j} \|x - y\|^2$$

$$D_{\text{means}}(C_i, C_j) = \|\mu_i - \mu_j\|^2$$



Example: microarray expression

- Measure gene expression
- Various experimental conditions
 - Cancer, normal
 - Time
 - Subjects
- Explore similarities
 - What genes change together?
 - What conditions are similar?
- Cluster on both genes and conditions



K-Means Clustering

- A simple clustering algorithm
- Iterate between
 - Updating the assignment of data to clusters
 - Updating the cluster's summarization
- Suppose we have K clusters, $c=1..K$
 - Represent clusters by locations μ_c
 - Example i has features x_i
 - Represent assignment of i^{th} example $z_i \in 1..K$
- Iterate until convergence:
 - For each datum, find the closest cluster

$$z_i = \arg \min_c \|x_i - \mu_c\|^2 \quad \forall i$$

- Set each cluster to the mean of all assigned data:

$$\forall c, \quad \mu_c = \frac{1}{N_c} \sum_{i \in S_c} x_i \quad S_c = \{i : z_i = c\}, \quad N_c = |S_c|$$

K-Means as optimization

- Optimizing the cost function

$$C(\underline{z}, \underline{\mu}) = \sum_i \|x_i - \mu_{z_i}\|^2$$

- Greedy descent technique
 - Steps
 - Choose closest cluster
 - Choose mean of assigned data
 - Each step only decreases the cost (why?)
- As with any descent method, beware of local minima
 - Algorithm behavior depends significantly on initialization
 - Many heuristics
 - Random (not bad); Farthest (sensitive); some mix maybe?

Choosing the number of clusters

- With cost function


$$C(\underline{z}, \underline{\mu}) = \sum_i \|x_i - \mu_{z_i}\|^2$$

what is the optimal value of k?

- Can increasing k ever increase the cost?
- This is a model complexity issue
 - Much like choosing lots of features – they only (seem to) help
 - But we want our clustering to *generalize* to new data
- One solution is to penalize for complexity
 - Bayesian information criterion (BIC)
 - Add (# parameters) * log(N) to the cost
 - Now more clusters can increase cost, if they don't help “enough”


Mixtures of Gaussians

- K-means algorithm
 - Assigned each example to exactly one cluster
 - What if clusters are overlapping?
 - Hard to tell which cluster is right
 - Maybe we should try to remain uncertain
 - Used Euclidean distance
 - What if cluster has a non-circular shape?
- Gaussian mixture models
 - Clusters modeled as Gaussians
 - Not just by their mean
 - EM algorithm: assign data to cluster with some *probability*



Multivariate Gaussian models

$$\mathcal{N}(\underline{x} ; \underline{\mu}, \Sigma) = \frac{1}{(2\pi)^{d/2}} |\Sigma|^{-1/2} \exp \left\{ -\frac{1}{2} (\underline{x} - \underline{\mu})^T \Sigma^{-1} (\underline{x} - \underline{\mu}) \right\}$$



Maximum Likelihood estimates

$$\hat{\mu} = \frac{1}{N} \sum_i x^{(i)}$$

$$\hat{\Sigma} = \frac{1}{N} \sum_i (x^{(i)} - \hat{\mu})^T (x^{(i)} - \hat{\mu})$$

**We'll model each cluster
using one of these Gaussian
"bells"...**

EM Algorithm: E-step

- Start with parameters describing each cluster
- Mean μ_c , Covariance Σ_c , “size” π_c
- E-step (“Expectation”)
 - For each datum (example) x_i ,
 - Compute “ r_{ic} ”, the probability that it belongs to cluster c
 - Compute its probability under model c
 - Normalize to sum to one (over clusters c)

$$r_{ic} = \frac{\pi_c \mathcal{N}(x_i ; \mu_c, \Sigma_c)}{\sum_{c'} \pi_{c'} \mathcal{N}(x_i ; \mu_{c'}, \Sigma_{c'})}$$

- If x_i is very likely under the c^{th} Gaussian, it gets high weight
- Denominator just makes r 's sum to one

EM Algorithm: M-step

- Start with assignment probabilities r_{ic}
- Update parameters: mean μ_c , Covariance Σ_c , “size” π_c
- M-step (“Maximization”)
 - For each cluster (Gaussian) x_c ,
 - Update its parameters using the (weighted) data points

$$N_c = \sum_i r_{ic} \quad \text{Total responsibility allocated to cluster c}$$

$$\pi_c = \frac{N_c}{N} \quad \text{Fraction of total assigned to cluster c}$$

$$\mu_c = \frac{1}{N_c} \sum_i r_{ic} x_i$$

Weighted mean of assigned data

$$\Sigma_c = \frac{1}{N_c} \sum_i r_{ic} (x_i - \mu_c)^T (x_i - \mu_c)$$

**Weighted covariance of assigned data
(use new weighted means here)**

Expectation-Maximization


- Each step increases the log-likelihood of our model

$$\log p(\underline{X}) = \sum_i \log \left[\sum_c \pi_c \mathcal{N}(x_i ; \mu_c, \Sigma_c) \right]$$

(we won't derive this, though)


- Iterate until convergence
 - Convergence guaranteed – another ascent method
- What should we do
 - If we want to choose a single cluster for an “answer”?
 - With new data we didn't see during training?

ANEMIA PATIENTS AND CONTROLS



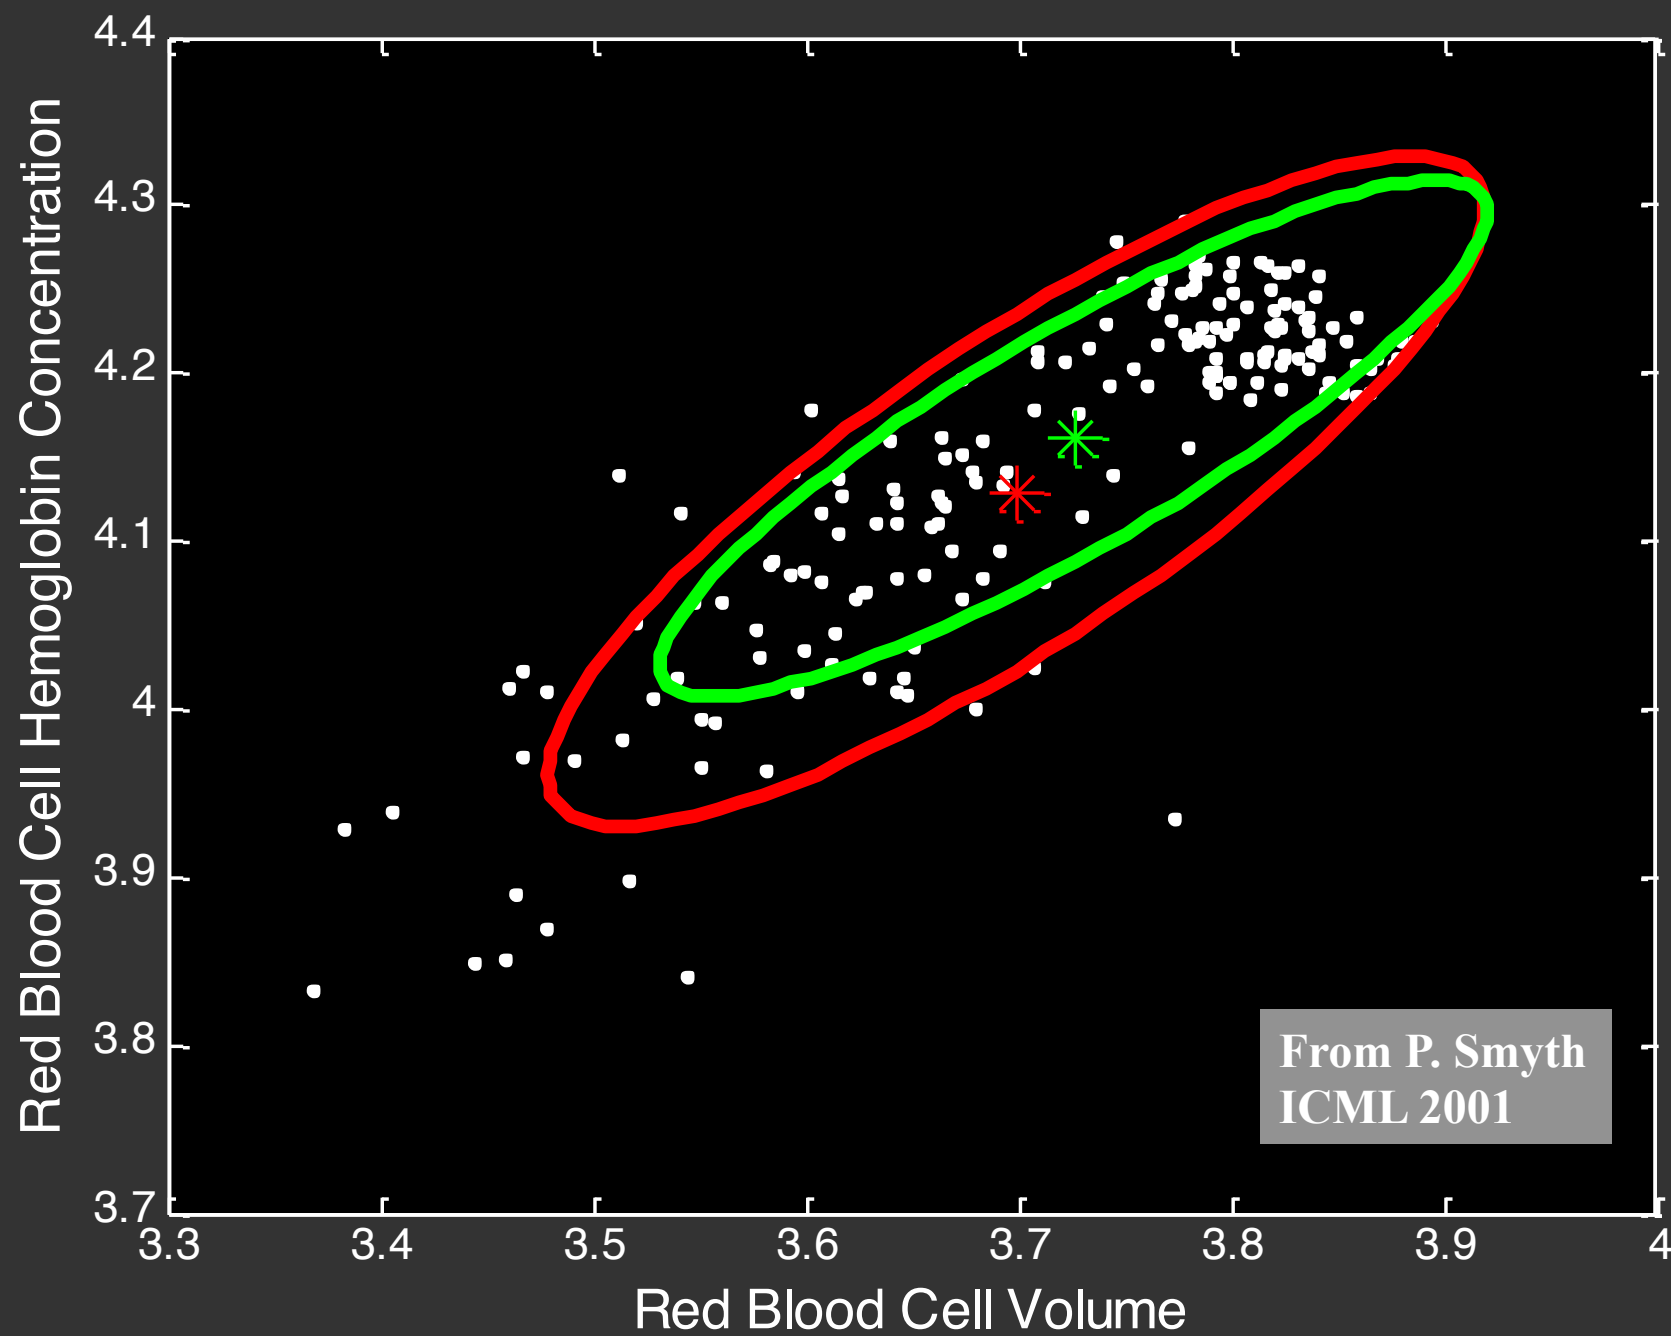
From P. Smyth
ICML 2001

EM ITERATION 1




From P. Smyth
ICML 2001


EM ITERATION 3



EM ITERATION 5




EM ITERATION 10




From P. Smyth
ICML 2001

EM ITERATION 15




EM ITERATION 25



From P. Smyth
ICML 2001

LOG-LIKELIHOOD AS A FUNCTION OF EM ITERATIONS



From P. Smyth
ICML 2001

EM and missing data

- EM is a general framework for partially observed data
 - “Complete data” x_i, z_i – features and assignments
 - Assignments z_i are missing (unobserved)
- EM corresponds to
 - Computing the distribution over all z_i given the parameters
 - Maximizing the “expected complete” log likelihood
 - GMMs = plug in “soft assignments”, but not always so easy
- Alternatives: Hard EM, Stochastic EM
 - Instead of expectations, just sample the z_i or choose best (often easier)
 - Called “imputing” the values of z_i
 - Hard EM: similar to EM, but less “smooth”
 - Stochastic EM: similar to EM, but with extra randomness
 - Not obvious when it has converged

Gibbs sampling for clustering

- Another technique for inferring uncertain cluster assignments
 - K-means: take the best assignment
 - EM: assign “partially”
 - Stochastic EM: sample assignment
 - All: choose best cluster descriptions given assignments
- Gibbs sampling (“Markov chain Monte Carlo”)
 - Assign randomly, probability equal to EM’s weight
 - *Sample* a cluster description given assignment
 - Requires a probability model over cluster parameters
- This doesn’t really find the “best” clustering
 - It eventually samples almost all “good” clusterings
 - Converges “in probability”, randomness helps us explore configurations
 - Also tells us about uncertainty of clustering
 - Disadvantage: not obvious when “done”

“Infinite” mixture models

- How many clusters are there?
- Gibbs sampling has an interesting solution
 - Write a distribution over k , the # of clusters
 - Sample k also
- Can do our sampling sequentially
 - Draw each z_i given all the others
 - Instead of sampling cluster parameters, marginalize them
 - Defines a distribution over groupings of data
- Now, for each z_i , sample
 - Join an existing cluster? Or, join a new cluster?
- What are these probabilities?
 - “Dirichlet process” mixture models

Parametric and Nonparametric Models

- Every model has some parameters
 - “The stuff you have to store to make your prediction”
 - Logistic regression: weights
 - Decision tree: feature to split, value at each level
 - Gaussian mixture model: means, covariances, sizes
- Parametric vs Nonparametric models
 - Parametric: fixed # of parameters
 - Nonparametric: # of parameters grows with more data
- What type are
 - Logistic regression?
 - Nearest neighbor prediction?
 - Decision trees?
 - Decision trees of depth < 3 ?
 - Gaussian mixture model?

Summary

- Clustering algorithms
 - Agglomerative clustering
 - K-means
 - Expectation-Maximization
- Open questions for each application
- What does it mean to be “close” or “similar”?
 - Depends on your particular problem...
- “Local” versus “global” notions of similarity
 - Former is easy, but we usually want the latter...
- Is it better to “understand” the data itself (unsupervised learning), to focus just on the final task (supervised learning), or both?